

# What is a Database Management System?

- A Database Management System (DBMS) is a software package designed to store and manage databases:
  1. Manages very large amounts of data.
  2. Supports efficient access to very large amounts of data.
  3. Supports concurrent access to very large amounts of data.
    - Example: bank and its ATM machines.
  4. Supports secure, atomic access to very large amounts of data.
    - Contrast two people editing the same UNIX file – last to write “wins” – with the problem if two people deduct money from the same account via ATM machines at the same time – new balance is wrong whichever writes last.

# Example: Online Bookseller

- Data = information on books (including categories, bestsellers, etc.), customers, pending orders, order histories, trends and preferences, etc.
  - Massive: many gigabytes at a minimum for medium-size bookseller, more if keep all order histories over all time, even more if keep images of book covers and sample pages
    - => Far too big for memory
  - Persistent: data outlives programs that operate on it
  - Multi-user: many people/programs accessing same database, or even same data, simultaneously
    - => Need careful controls



# Files vs. DBMS

- Application must stage large datasets between main memory and secondary storage (e.g., buffering, page-oriented access, 32-bit addressing, etc.)
- Special code for different queries
- Must protect data from inconsistency due to multiple concurrent users
- Crash recovery
- Security and access control

# What is a Relational Database?

- Based on the relational model (tables):

acct #	name	balance
12345	Sally	1000.21
34567	Sue	285.48
...	...	...

- Today used in *most* DBMS's.

# The DBMS Marketplace

- Relational DBMS companies – Oracle, Sybase – are among the largest software companies in the world.
- IBM offers its relational DB2 system. With IMS, a nonrelational system, IBM is by some accounts the largest DBMS vendor in the world.
- Microsoft offers SQL-Server, plus Microsoft Access for the cheap DBMS on the desktop, answered by “lite” systems from other competitors.
- Relational companies also challenged by “object-oriented DB” companies.
- But countered with “object-relational” systems, which retain the relational core while allowing type extension as in OO systems.

# Three Aspects to Studying DBMS's

1. Modeling and design of databases.
  - Allows exploration of issues before committing to an implementation.
2. Programming: queries and DB operations like update.
  - SQL = “intergalactic dataspeak.”
3. DBMS implementation.

# Query Languages

Employee	
Name	Dept

Department	
Dept	Manager

SQL

```
SELECT Manager
FROM Employee, Department
WHERE Employee.name = "Clark Kent"
      AND Employee.Dept = Department.Dept
```

Query Language

Data definition language (DDL) ~ like type defs in C or Pascal

Data Manipulation Language (DML)

Query (SELECT)

UPDATE <relation name>

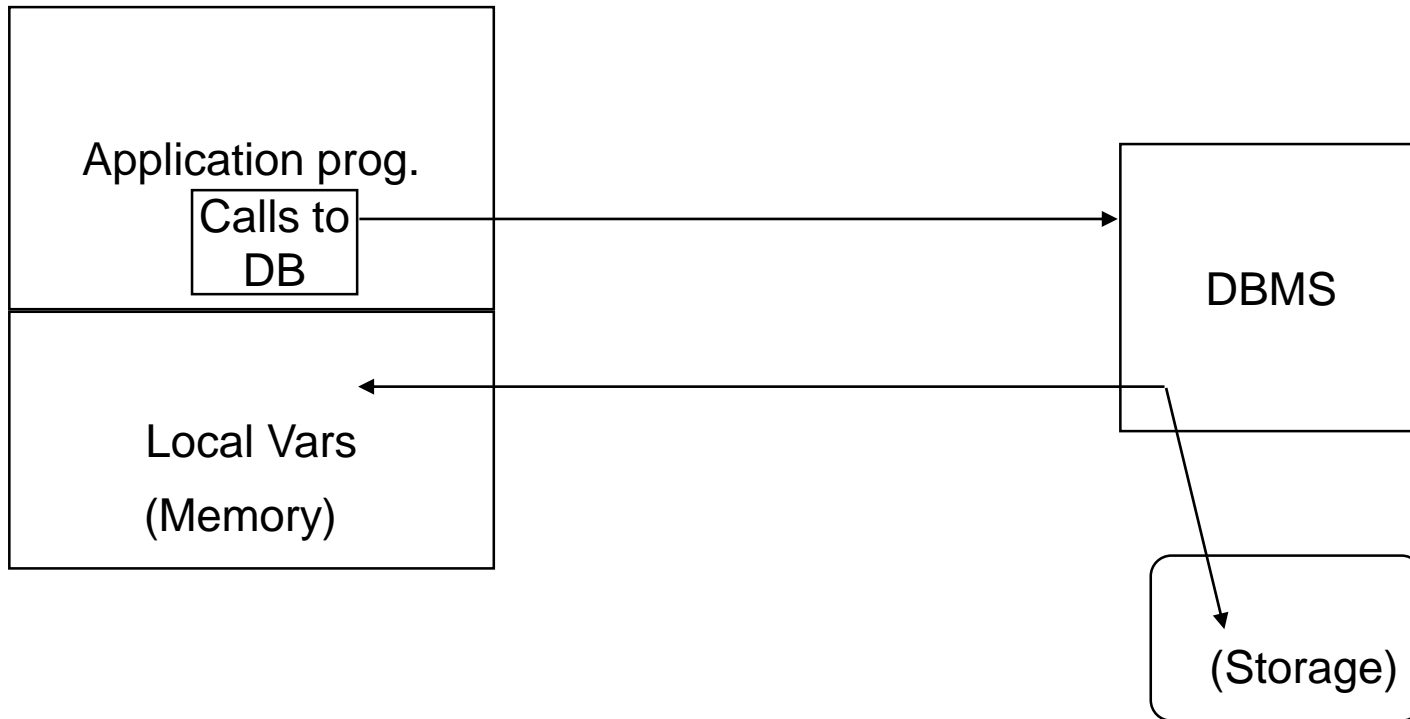
SET <attribute> = <new-value>

WHERE <condition>



# Host Languages

C, C++, Java, Lisp, COBOL



- Host language is completely general (Turing complete)
- Query language—less general "non procedural" and optimizable

# Relational Model

Relational model is good for:

- Large amounts of data —> simple operations
- Navigate among small number of relations

Difficult Applications for relational model:

- VLSI Design (CAD in general)
- CASE
- Graphical Data

# Other Models

Where number of "relations" is large, relationships are complex

- Object Data Model
- Logic Data Model

## OBJECT DATA MODEL

1. Complex Objects – Nested Structure (pointers or references)
2. Encapsulation, set of Methods/Access functions
3. Object Identity
4. Inheritance – Defining new classes like old classes

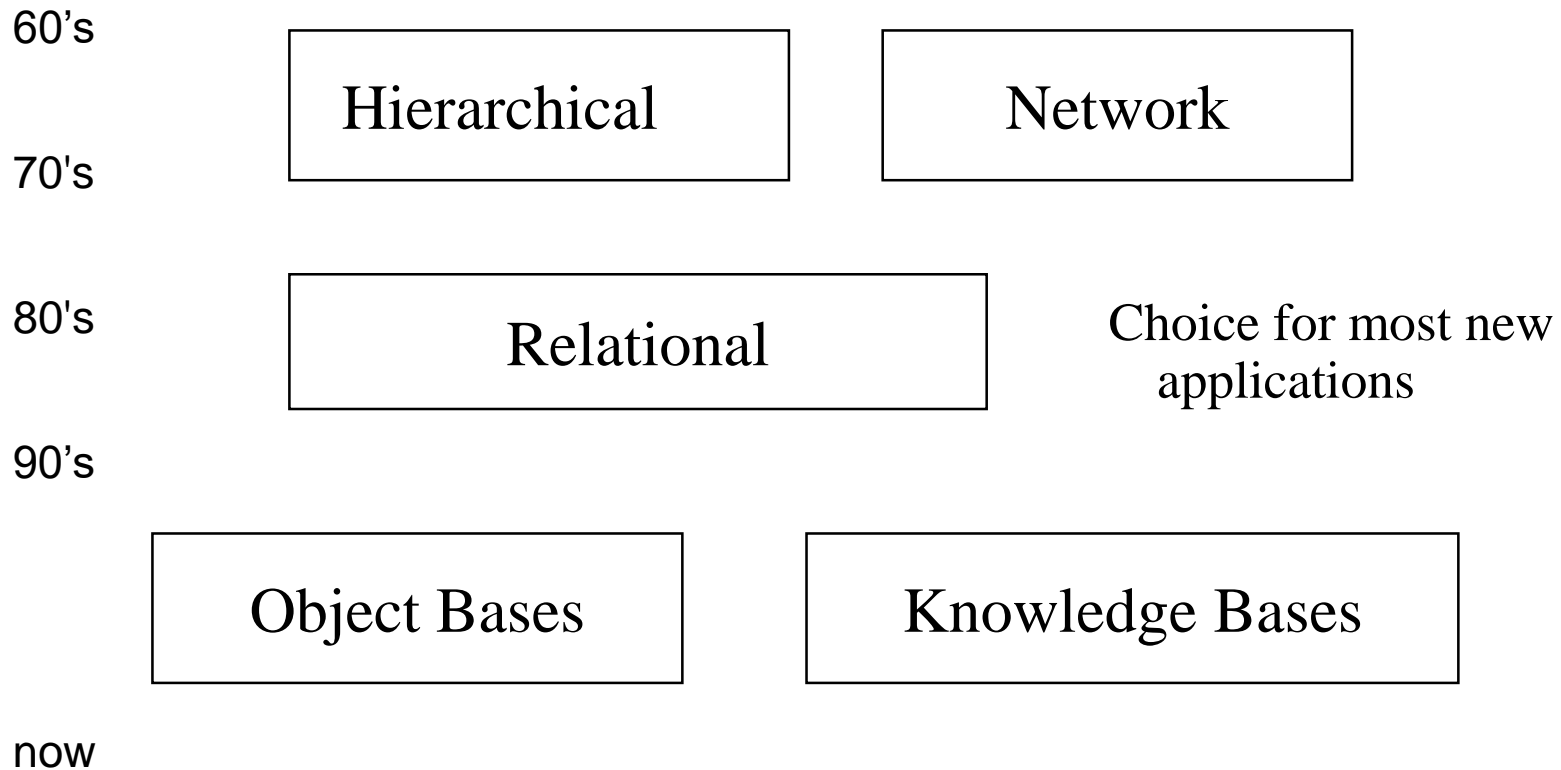
Object model: usually find objects via explicit navigation  
Also query language in some systems

# Other Models

## LOGIC (Horn Clause) DATA MODEL

- Prolog, Datalog:  
if A1 and A2 then B  
B:- A1 and A2
- Functions  $s(5) = 6$  (successor)
- Predicates with Arguments:  $\text{sum}(X,Y,Z) \leftarrow X + Y = Z$   
 $\text{sum}(X,0,X)$  means  $X + 0 = X$  (always true for all X)  
 $\text{sum}(X,s(Y),s(Z)):-\text{sum}(X,Y,Z)$  means  $X+(Y+1) = (Z+1)$  if  $X + Y = Z$
- More powerful than relational  
Can Compute Transitive Closure  
 $\text{edge}(X,Y).$   
 $\text{path}(X,Y) :- \text{edge}(X,Y).$   
 $\text{path}(X,Z) :- \text{path}(X,Y) \ \& \ \text{edge}(Y,Z).$

# Data Models



# Why Use a DBMS?

- Data independence and efficient access.
- Reduced application development time.
- Data integrity and security.
- Uniform data administration.
- Concurrent access, recovery from crashes.

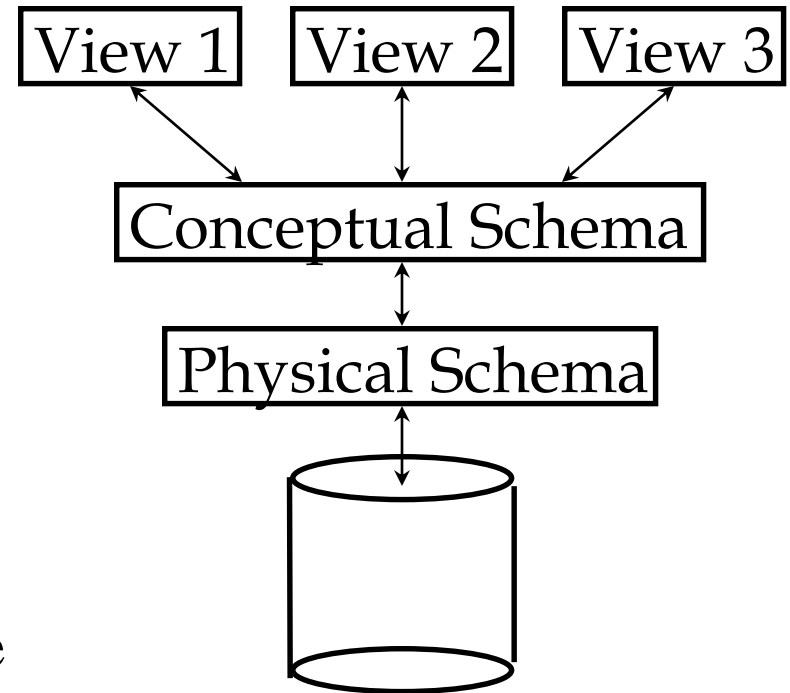
# Data Independence \*

- Applications insulated from how data is structured and stored.
- Logical data independence: Protection from changes in *logical* structure of data.
- Physical data independence: Protection from changes in *physical* structure of data.

✉ *One of the most important benefits of using a DBMS!*

# Levels of Abstraction

- Many views, single conceptual (logical) schema and physical schema.
  - Views describe how users see the data.
  - Conceptual schema defines logical structure
  - Physical schema describes the files and indexes used.



✉ Schemas are defined using DDL; data is modified/queried using DML.



# Concurrency Control

- Concurrent execution of user programs is essential for good DBMS performance.
  - Because disk accesses are frequent, and relatively slow, it is important to keep the cpu humming by working on several user programs concurrently.
- Interleaving actions of different user programs can lead to inconsistency: e.g., cheque is cleared while account balance is being computed.
- DBMS ensures such problems don't arise: users can pretend they are using a single-user system.

# Transaction: An Execution of a DB Program

- Key concept is transaction, which is an *atomic* sequence of database actions (reads/writes).
- Each transaction, executed completely, must leave the DB in a consistent state if DB is consistent when the transaction begins.
  - Users can specify some simple integrity constraints on the data, and the DBMS will enforce these constraints.
  - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
  - Thus, ensuring that a transaction (run alone) preserves consistency is ultimately the user's responsibility!

# Scheduling Concurrent Transactions

- DBMS ensures that execution of  $\{T_1, \dots, T_n\}$  is equivalent to some serial execution  $T_1' \dots T_n'$ .
  - Before reading/writing an object, a transaction requests a lock on the object, and waits till the DBMS gives it the lock. All locks are released at the end of the transaction. (Strict 2PL locking protocol.)
  - **Idea:** If an action of  $T_i$  (say, writing  $X$ ) affects  $T_j$  (which perhaps reads  $X$ ), one of them, say  $T_i$ , will obtain the lock on  $X$  first and  $T_j$  is forced to wait until  $T_i$  completes; this effectively orders the transactions.
  - What if  $T_j$  already has a lock on  $Y$  and  $T_i$  later requests a lock on  $Y$ ? (Deadlock!)  $T_i$  or  $T_j$  is aborted and restarted!

# Ensuring Atomicity

- DBMS ensures *atomicity* (all-or-nothing property) even if system crashes in the middle of a Xact.
- **Idea:** Keep a log (history) of all actions carried out by the DBMS while executing a set of Xacts:
  - **Before** a change is made to the database, the corresponding log entry is forced to a safe location. (WAL protocol; OS support for this is often inadequate.)
  - After a crash, the effects of partially executed transactions are undone using the log. (Thanks to WAL, if log entry wasn't saved before the crash, corresponding change was not applied to database!)

# The Log

- The following actions are recorded in the log:
  - *Ti writes an object*: The old value and the new value.
    - Log record must go to disk before the changed page!
  - *Ti commits/aborts*: A log record indicating this action.
- Log records chained together by Xact id, so it's easy to undo a specific Xact (e.g., to resolve a deadlock).
- Log is often *duplexed* and *archived* on “stable” storage.
- All log related activities (and in fact, all CC related activities such as lock/unlock, dealing with deadlocks etc.) are handled transparently by the DBMS.

# Databases make these folks happy

- End users and DBMS vendors
- DB application programmers
  - e.g., smart webmasters
- Database administrator (DBA)
  - Designs logical /physical schemas
  - Handles security and authorization
  - Data availability, crash recovery
  - Database tuning as needs evolve

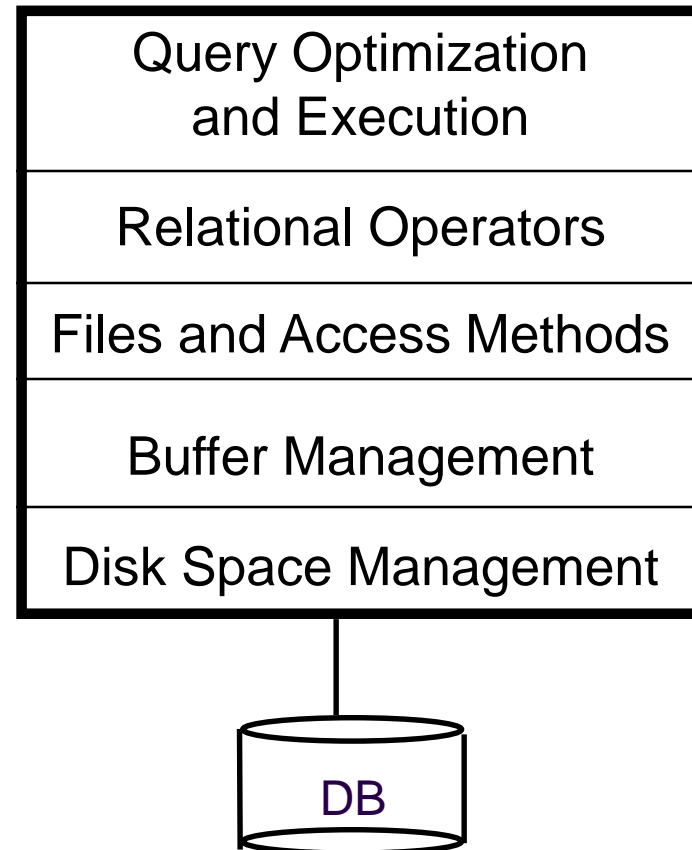


*Must understand how a DBMS works!*

# Structure of a DBMS

These layers must consider concurrency control and recovery

- A typical DBMS has a layered architecture.
- The figure does not show the concurrency control and recovery components.
- This is one of several possible architectures; each system has its own variations.



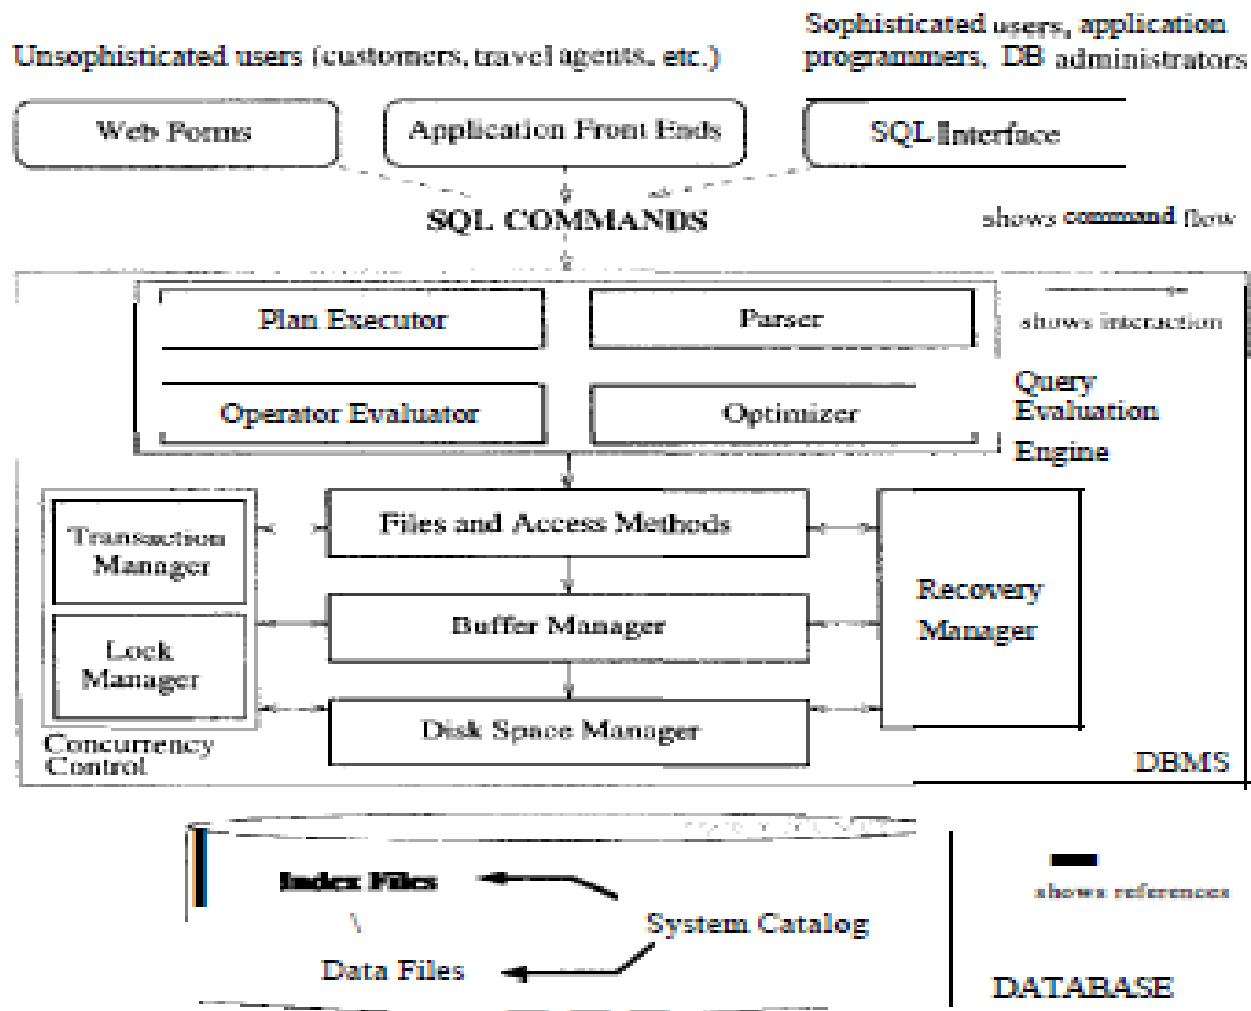


Figure 1.3 Architecture of a DBMS



# Summary

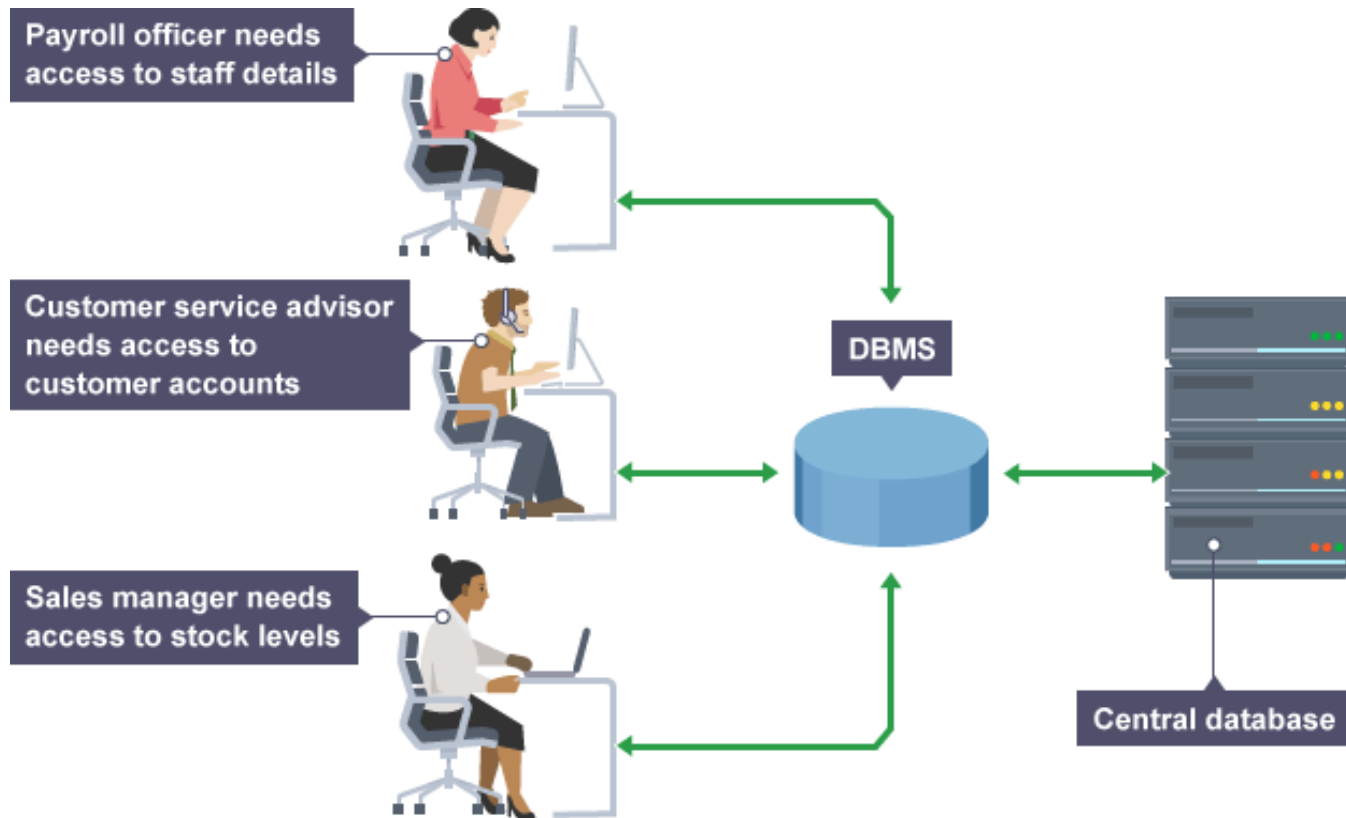
- DBMS used to maintain, query large datasets.
- Benefits include recovery from system crashes, concurrent access, quick application development, data integrity and security.
- Levels of abstraction give data independence.
- A DBMS typically has a layered architecture.
- DBAs hold responsible jobs are well-paid! ☺
- DBMS R&D is one of the broadest, most exciting areas in CS.



# **The Entity-Relationship Model**

# Data Model

- Move from informal description of what user wants to
- Precise description of what can be implemented in a DBMS



# Steps in developing a database

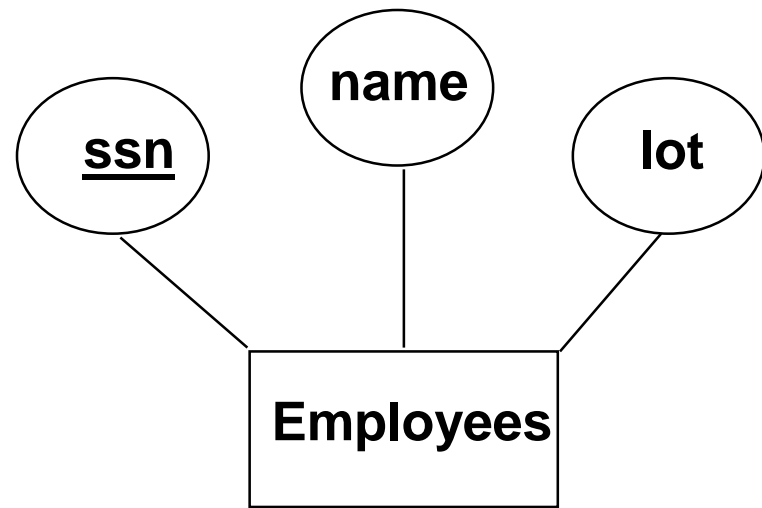
- Requirements analysis →
- Conceptual Database design →
- Logical Database design →
- Schema refinement →
- Physical Database design
- Applications and security

# Conceptual design: (ER Model is used at this stage.)

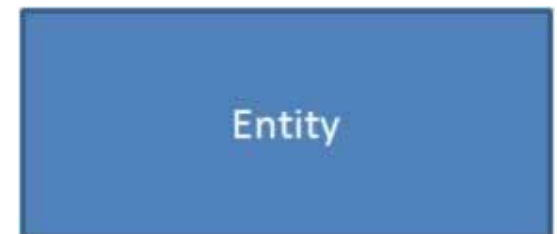
- What are the *entities* and *relationships* in the enterprise?
- What information about these entities and relationships should we store in the database?
- What are the *integrity constraints* or *business rules* that hold?
- A database 'schema' in the ER Model can be represented pictorially (*ER diagrams*).
- Can map an ER diagram into a relational schema.

- The ER diagram is just an approximate description of the data, constructed through a subjective evaluation of the information collected during requirements analysis.
- A more careful analysis can often refine the logical schema obtained at the end of Step 3

# ER Model Basics



- Entity: Real-world object distinguishable from other objects. An entity is described using a set of attributes. Each attribute has a *domain*.
- Entity Set: A collection of similar entities. E.g., all employees.
  - All entities in an entity set have the same set of attributes. (Until we consider ISA hierarchies, anyway!)
  - Each entity set has a *key*.





- Two classes of entities are there:
  - Regular Entity or Strong Entity
  - Weak Entity



## Attribute

Properties/characteristics which describe entities are called attributes.

## Key Attribute



## Multivalued Attributes

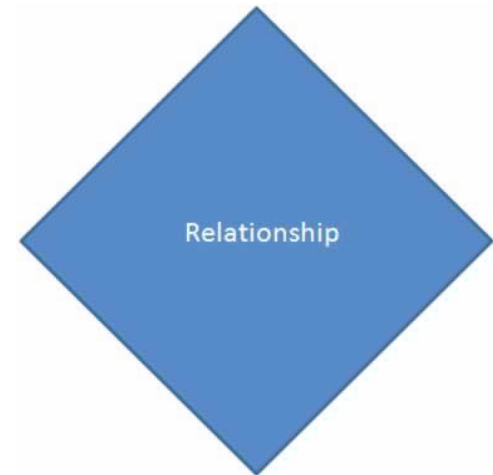


Derived Attribute



## Relationships

Associations between entities are called relationships



# Keys

- Minimal set of attributes which uniquely identify an instance of a entity
- Many candidate keys choose one to be a primary keys
- SSN vs Name ... key must be unique

# Keys

## Types of keys in DBMS

**Primary Key** – A primary is a column or set of columns in a table that uniquely identifies tuples (rows) in that table. (ex: In a student table SRN or studID)

**Super Key** – A super key is a set of one or more columns (attributes) to uniquely identify rows in a table. (ex: **Employee table (Emp\_SSN, Emp\_num, Emp\_name)**)

**Candidate Key** – A super key with no redundant attribute is known as candidate key, Candidate key is also called as minimal super key. Primary keys can be selected from Candidate keys.

**Alternate Key** – Out of all candidate keys, only one gets selected as primary key, remaining keys are known as alternate or secondary keys.

**Composite Key** – A key that consists of more than one attribute to uniquely identify rows (also known as records & tuples) in a table is called composite key.

**Foreign Key** – Foreign keys are the columns of a table that points to the primary key of another table. They act as a cross-reference between tables.

for better understanding of the key. Table “**Branch Info**”

	Branch_Id	Branch_Name	Branch_Code
	1	Computer Scie...	CSE
	2	Electronics	ECE
	3	Mechanical	MCE
	4	Information Te...	ITE
	5	Civil	CVE
▶*	NULL	NULL	NULL

*Table: Branch\_Info*

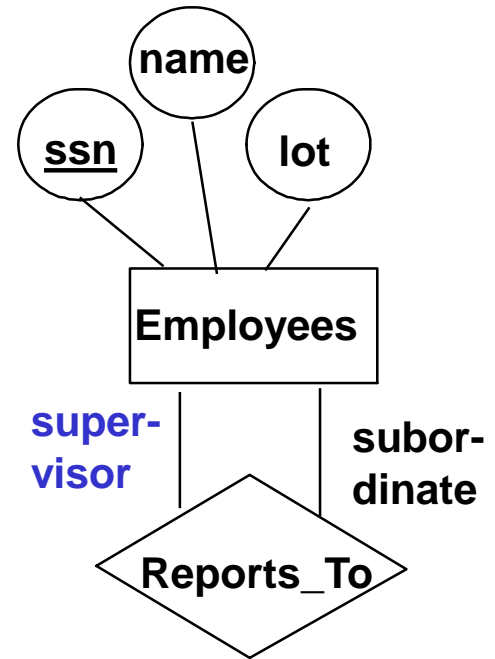
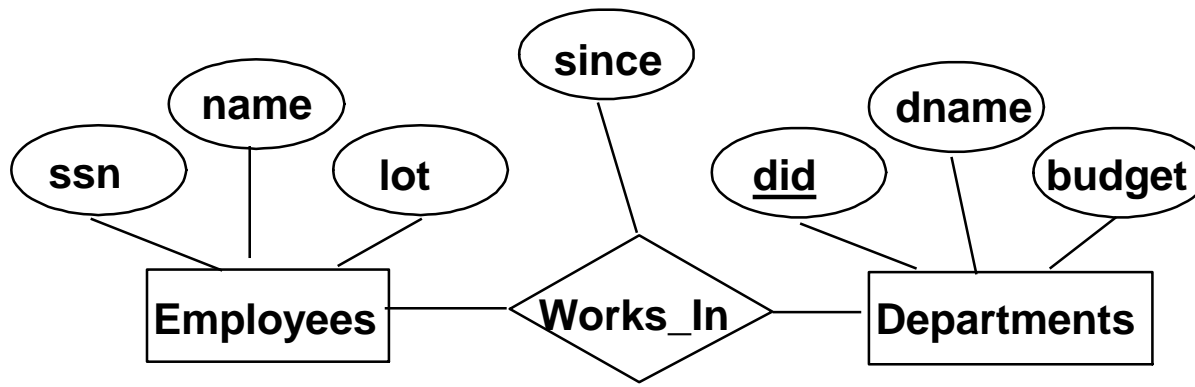
Possible Candidate Keys in Branch\_Info table.

- Branch\_Id
- Branch\_Name
- Branch\_Code

Among the above candidate keys one key can be primary key which uniquely identifies the entire relation. (Branch\_Id), Alternate Key (Secondary Key) is Branch\_name or Branch\_code.

- Unique Key

# ER Model Basics (Contd.)



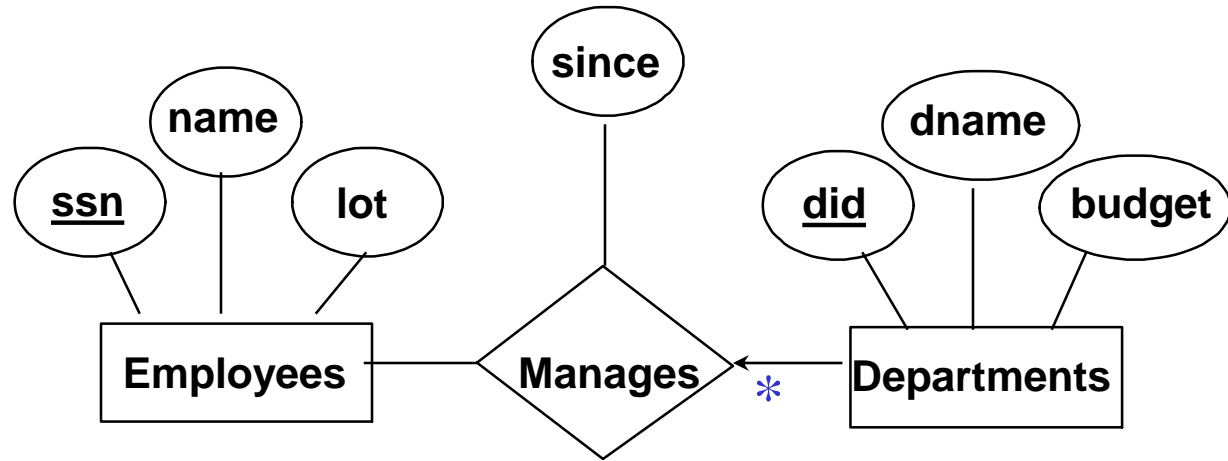
- **Relationship**: Association among two or more entities. E.g., Attishoo works in Pharmacy department.
- **Relationship Set**: Collection of similar relationships.
  - An n-ary relationship set  $R$  relates  $n$  entity sets  $E_1 \dots E_n$ ; each relationship in  $R$  involves entities  $e_1 \in E_1, \dots, e_n \in E_n$
  - Same entity set could participate in different relationship sets, or in different “**roles**” in same set.

# Key Constraints

- Consider Works\_In: An employee can work in many departments; a dept can have many employees.

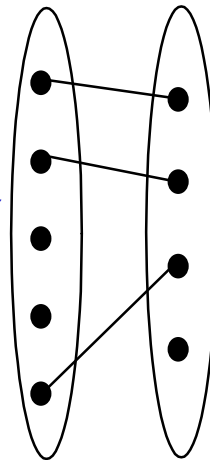
So many to many

# Key Constraints

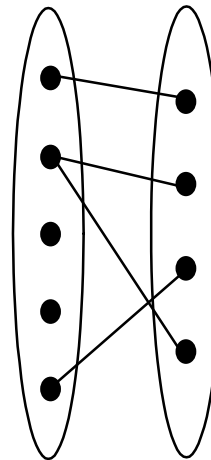


In contrast, each dept has at most one manager, according to the key constraint on Manages.

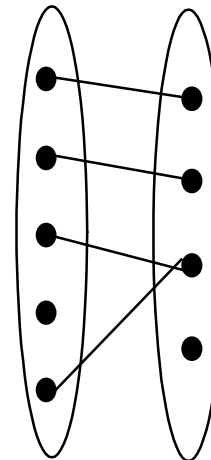
\* arrow indicates that given a dept it uniquely determines the Manages relationship in which it appears



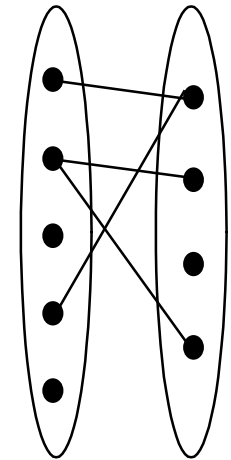
1-to-1



1-to Many



Many-to-1



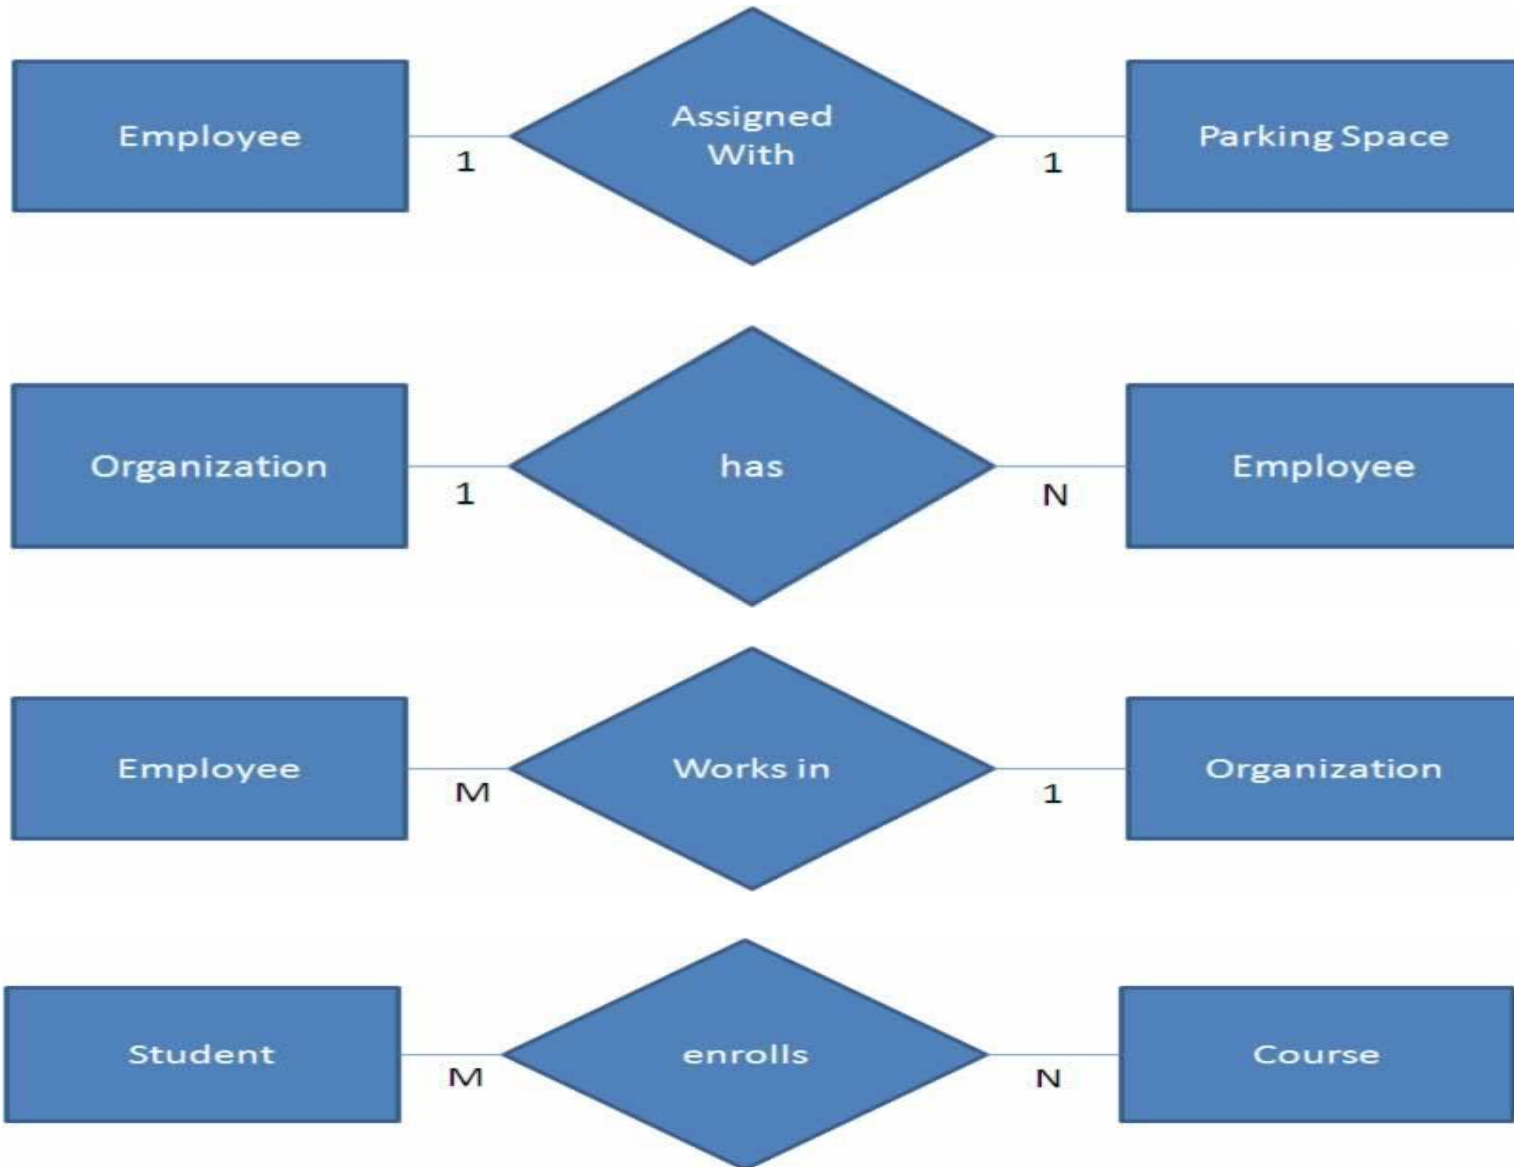
Many-to-Many



# Cardinality of a Relationship

- Relationship cardinalities specify how many of each entity type is allowed. Relationships can have four possible connectivity's as given below.
- 1. One to one (1:1) relationship
- 2. One to many (1:N) relationship
- 3. Many to one (M:1) relationship
- 4. Many to many (M:N) relationship

# Examples

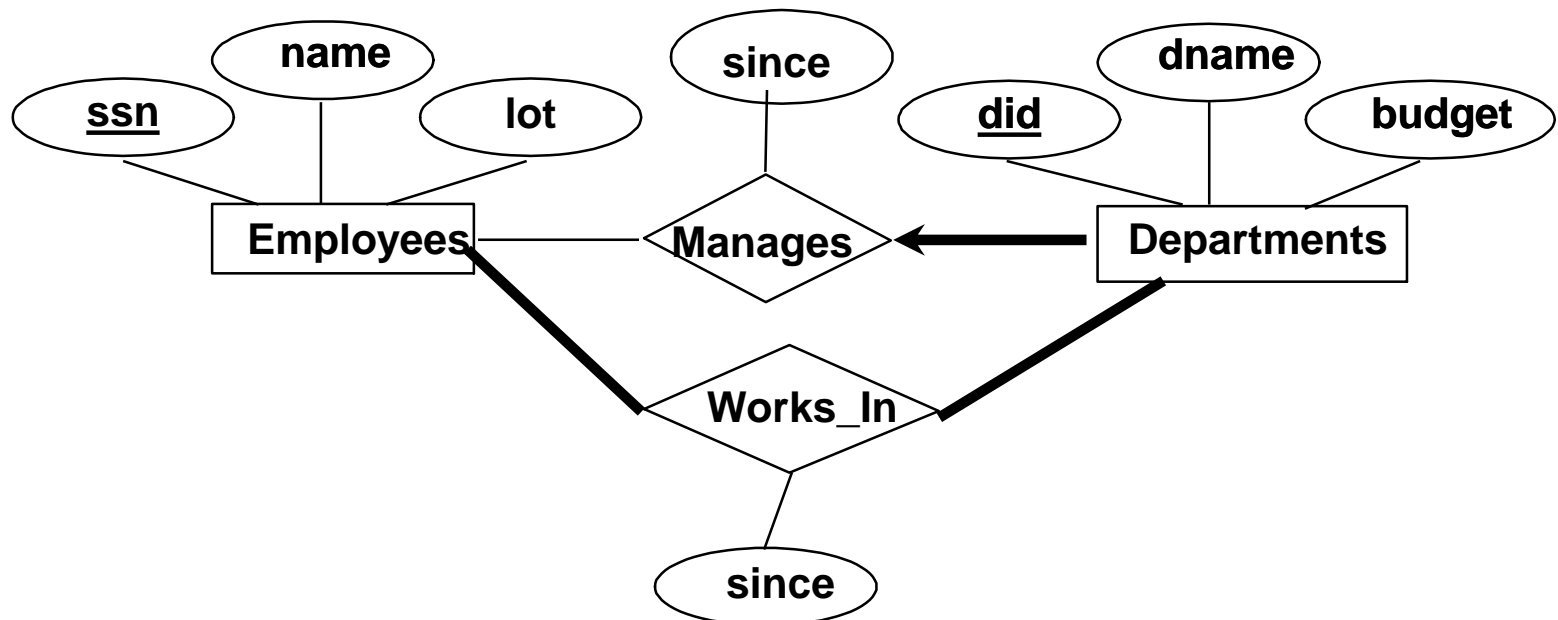


# Degree of a Relationship

- Degree of a relationship is the number of entity types involved. The n-ary relationship is the general form for degree n.
- Special cases are unary, binary, and ternary where the degree is 1, 2, and 3, respectively.
- Example for unary relationship : An employee is a manager of another employee
- Example for binary relationship : An employee works-for department.
- Example for ternary relationship : customer purchase item from a shop keeper

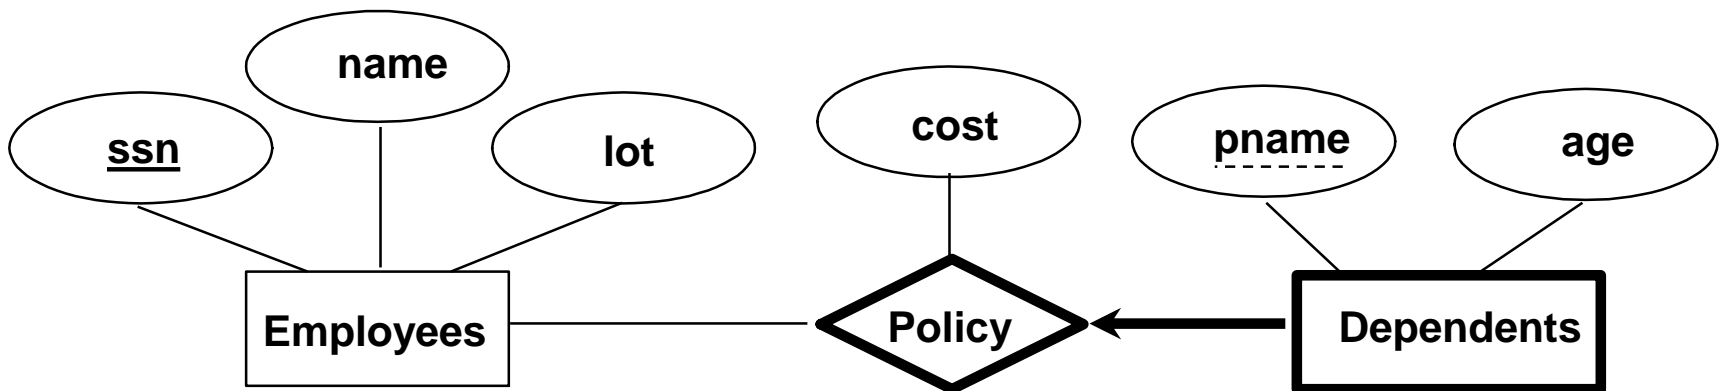
# Participation Constraints

- Does every department have a manager?
  - If so, this is a *participation constraint*: the participation of Departments in Manages is said to be *total* (vs. *partial*).
- Every Department entity must appear in an instance of the relationship Works\_In (have an employee) and every Employee must be in a Department
- Both Employees and Departments participate totally in Works\_In

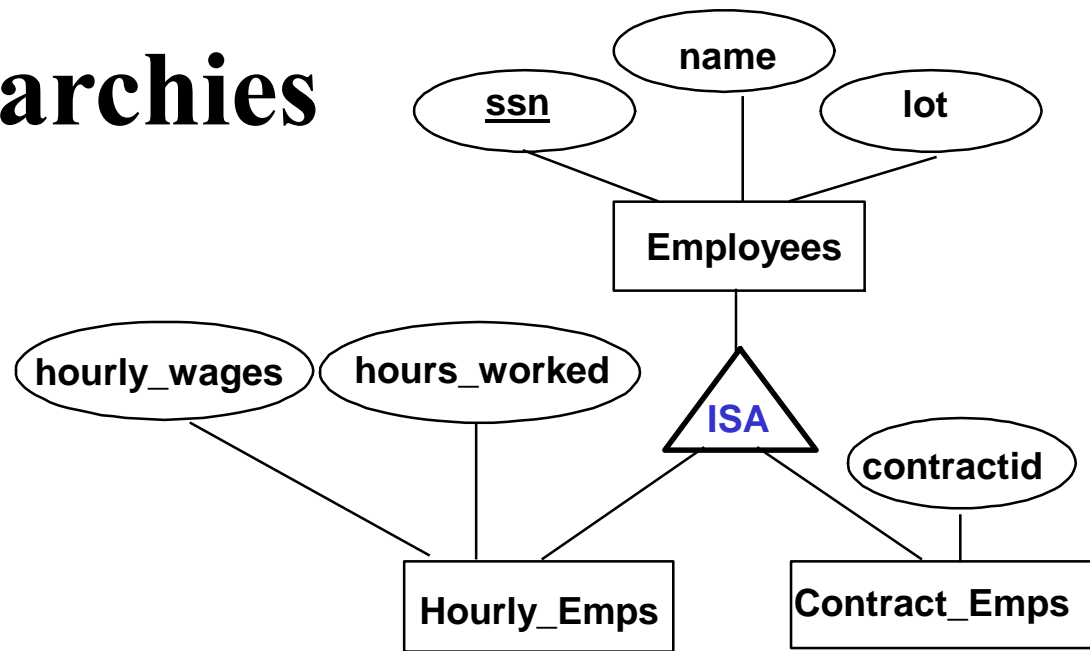


# Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
  - Weak entity set must have total participation in this *identifying* relationship set.



# ISA ('is a') Hierarchies

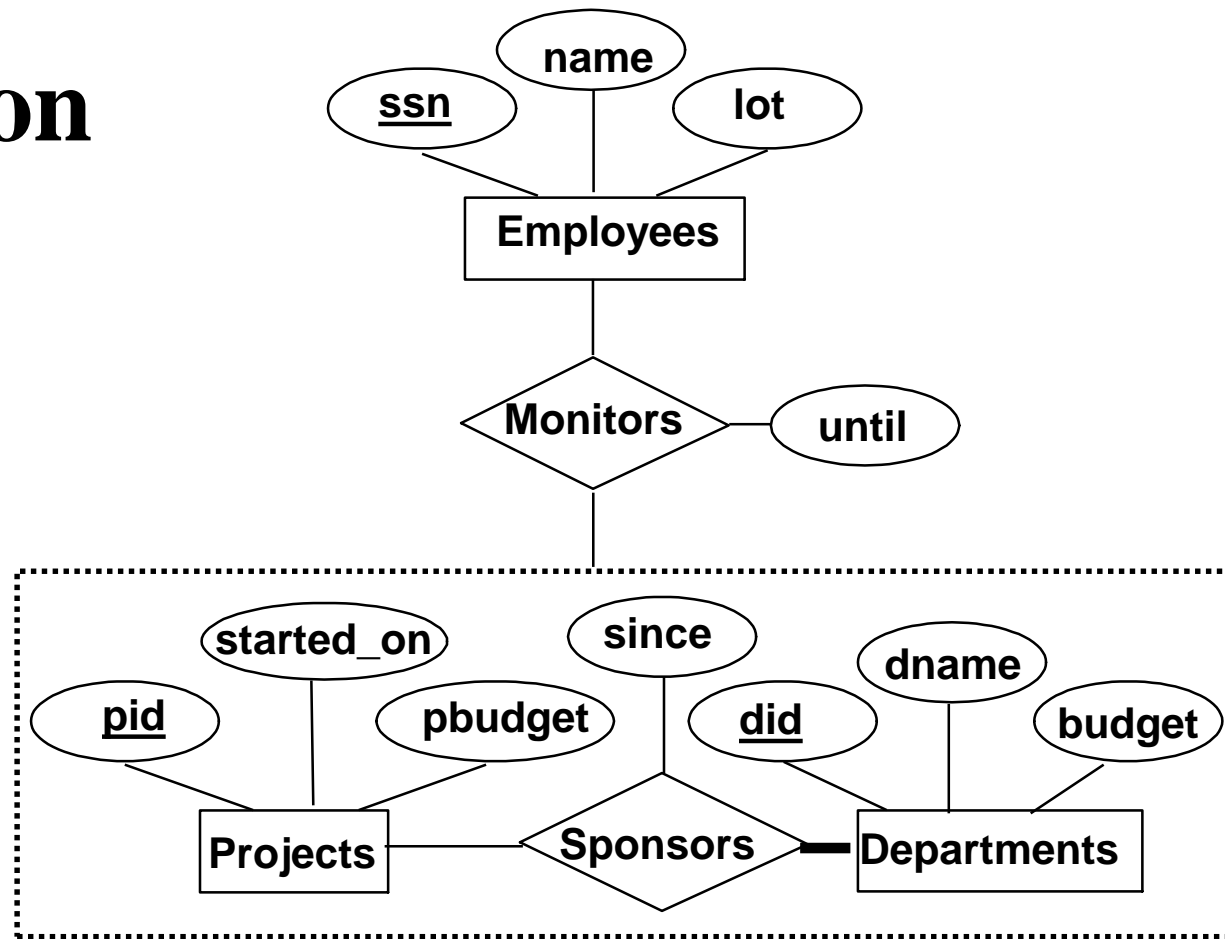


- ❖ As in C++, attributes can be inherited.
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.
- ❖ Upwards is generalization. Down is specialization

# Constraints in ISA relation

- *Overlap constraints:* Can Joe be an Hourly\_Emps as well as a Contract\_Emps entity? (*Allowed/disallowed*)
- *Covering constraints:* Does every Employees entity also have to be an Hourly\_Emps or a Contract\_Emps entity? (*Yes/no*)
- Reasons for using ISA:
  - To add descriptive attributes specific to a subclass.
  - To identify entities that participate in a relationship.

# Aggregation



- Used when we have to model a relationship involving (entity sets and) a *relationship set*.

Aggregation allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships.



# Aggregation vs. ternary relationship:

- Monitors in last example is a distinct relationship, with a descriptive attribute.
- Also, can say that each sponsorship is monitored by at most one employee.

# Conceptual Design Using the ER Model

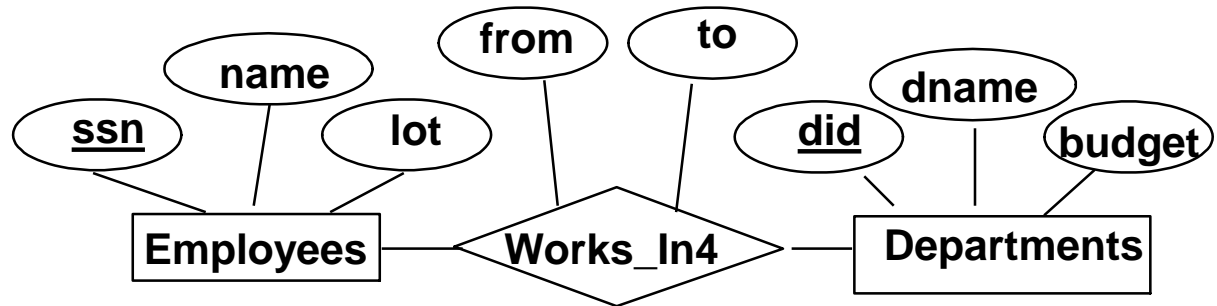
- Design choices:
  - Should a concept be modeled as an entity or a relationship?
  - Identifying relationships: Binary or ternary? Aggregation?
- Constraints in the ER Model:
  - A lot of data semantics can (and should) be captured.
  - But some constraints cannot be captured in ER diagrams.

# Entity vs. Attribute

- Should *address* be an attribute of Employees or an entity (connected to Employees by a relationship)?
- Depends upon the use we want to make of address information, and the semantics of the data:
  - If we have several addresses per employee, *address* must be an entity (since attributes cannot be set-valued).
  - If the structure (city, street, etc.) is important, e.g., we want to retrieve employees in a given city, *address* must be modeled as an entity (since attribute values are atomic).

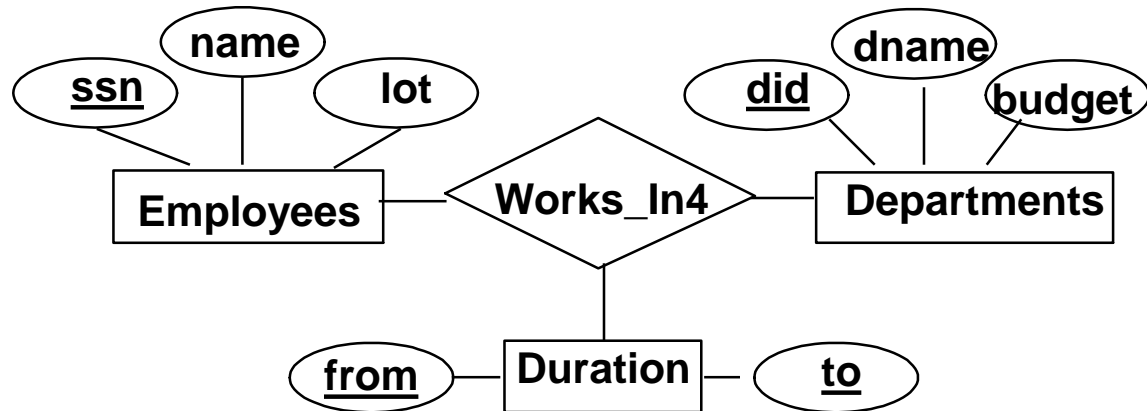
# Entity vs. Attribute (Contd.)

- Works\_In4 does not allow an employee to work in a department for two or more periods.

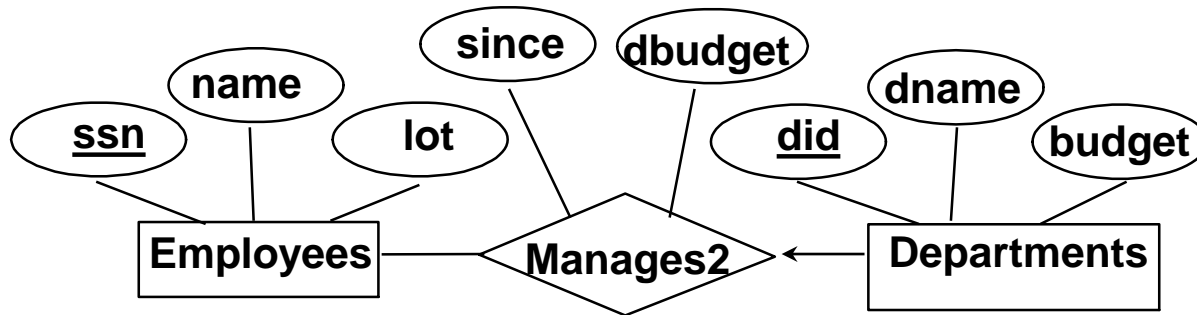


- Similar to the problem of wanting to record several addresses for an employee: We want to record *several values of the descriptive attributes for each instance of this relationship*.

Accomplished by introducing new entity set, **Duration**.

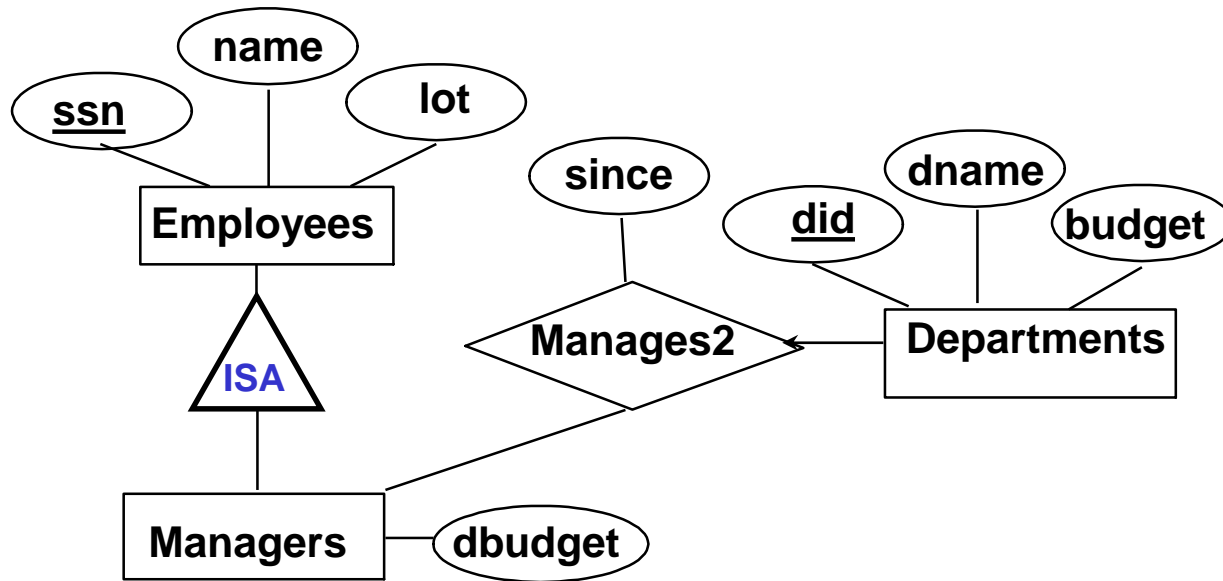


# Entity vs. Relationship

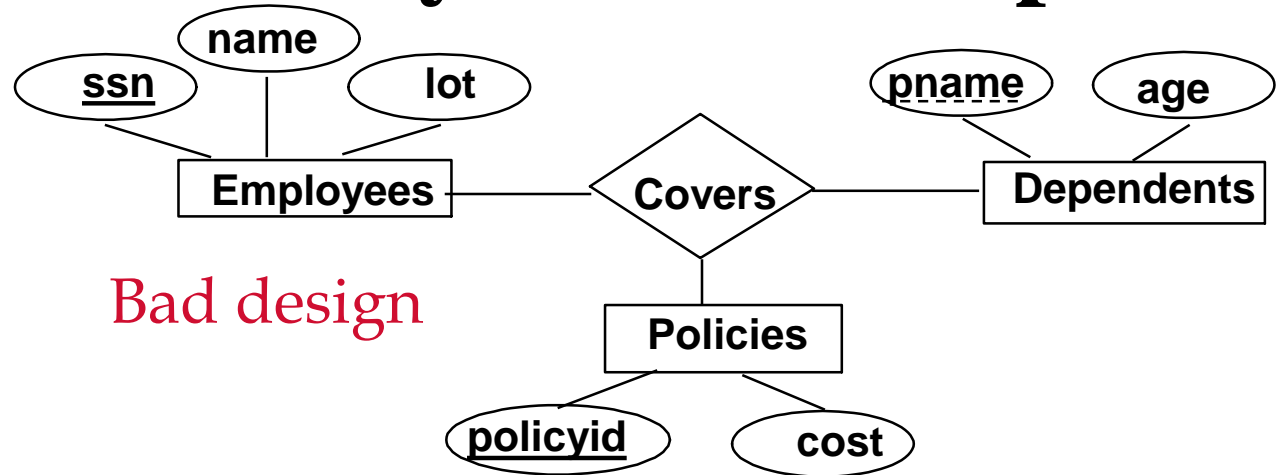


- ER diagram OK if a manager gets a separate discretionary budget for each dept.
- What if a manager gets a discretionary budget that covers *all* managed depts?
  - **Redundancy**: *dbudget* stored for each dept managed by manager.
  - **Misleading**: Suggests *dbudget* associated with department-mgr combination.

# This fixes the problem!

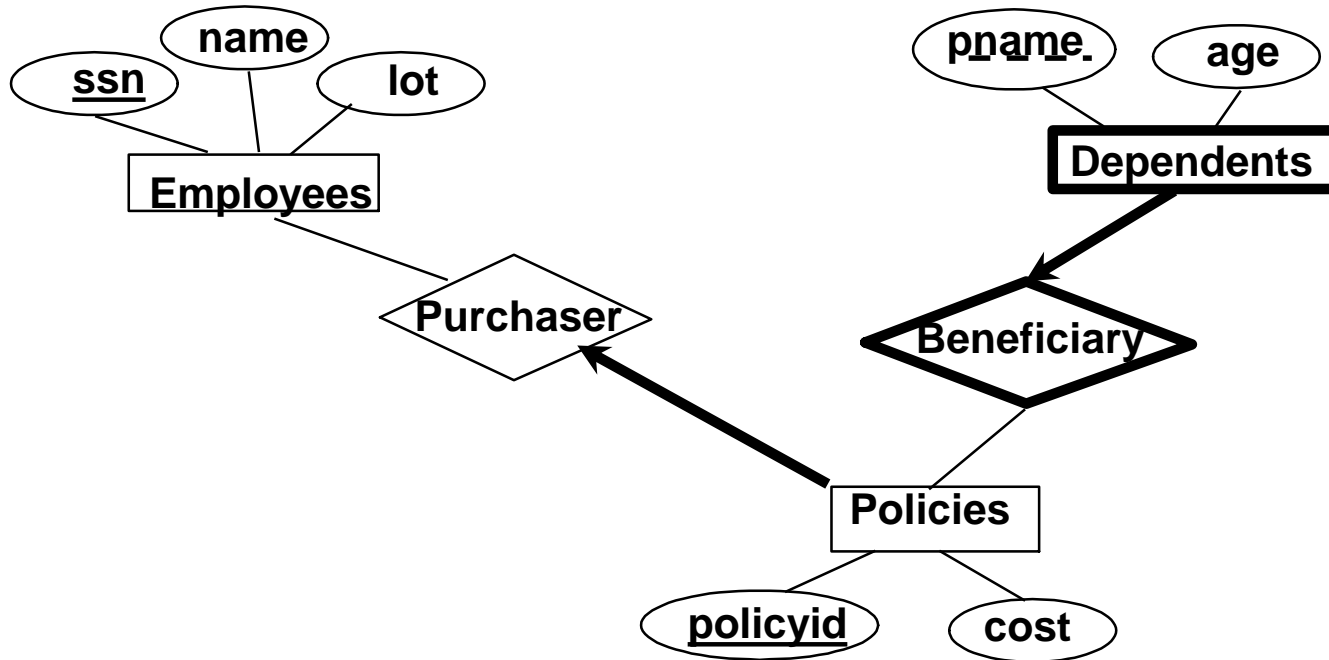


# Binary vs. Ternary Relationships



- If each policy is owned by just 1 employee, and each dependent is tied to the covering policy, first diagram is inaccurate.
- What are the additional constraints do we need?

# Better design



- Key constraint
- Total participation of policies in purchaser relationship



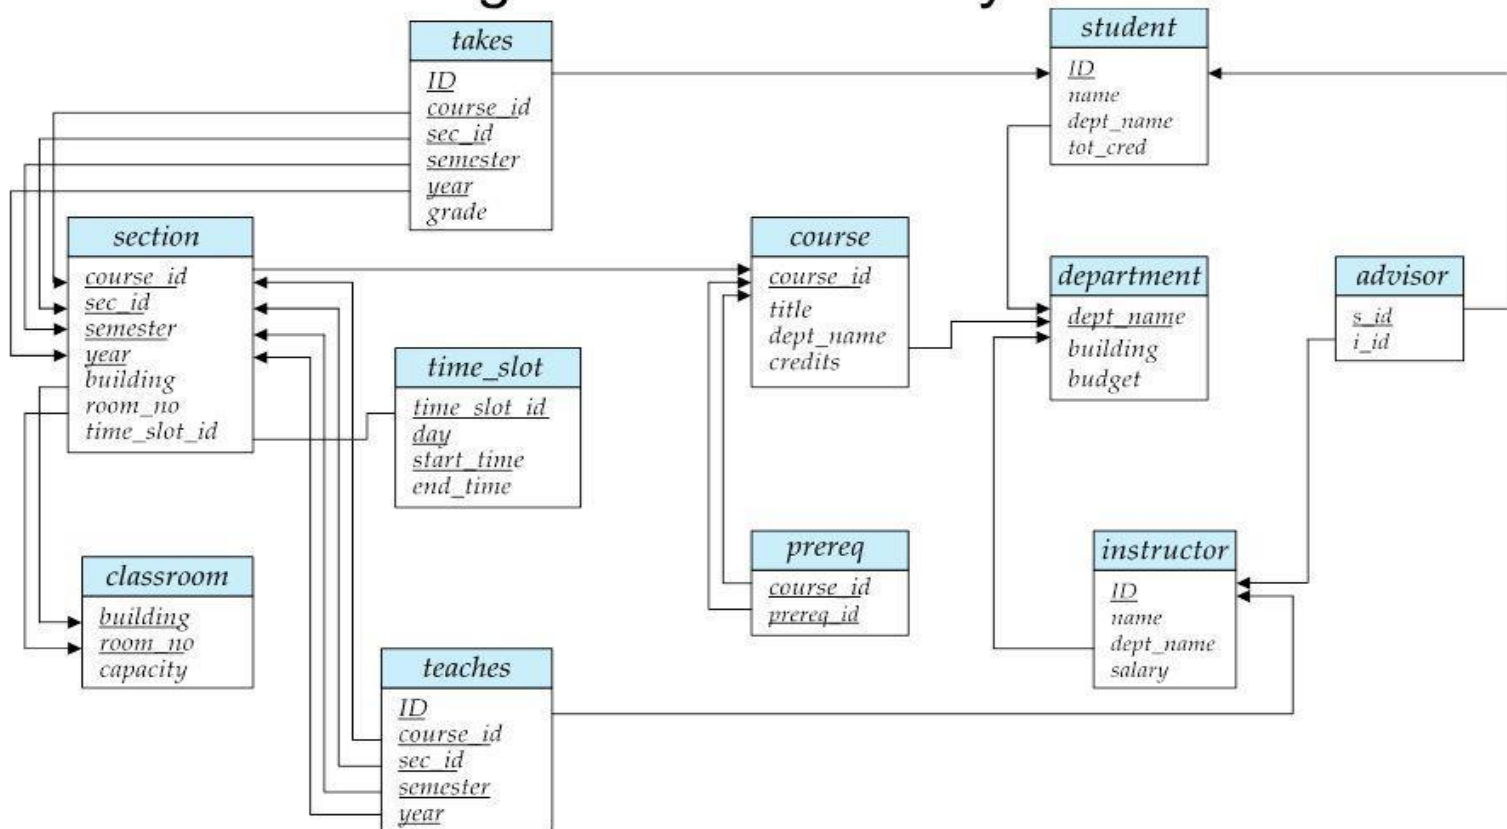
# Binary vs. Ternary Relationships (Contd.)

- Previous example illustrated a case when two binary relationships were better than one ternary relationship.
- An example in the other direction: a ternary relation **Contracts** relates entity sets **Parts**, **Departments** and **Suppliers**, and has descriptive attribute *qty*. No combination of binary relationships is an adequate substitute:
  - S “can-supply” P, D “needs” P, and D “deals-with” S does not imply that D has agreed to buy P from S.
  - How do we record *qty*?

# Summary of Conceptual Design

- *Conceptual design follows requirements analysis,*
  - Yields a high-level description of data to be stored
- ER model popular for conceptual design
  - Constructs are expressive, close to the way people think about their applications.
- Basic constructs: *entities, relationships, and attributes* (of entities and relationships).
- Some additional constructs: *weak entities, ISA hierarchies, and aggregation.*
- Note: There are many variations on ER model.

# Schema Diagram for University Database



# Summary of ER (Contd.)

- Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set.
  - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
  - Constraints play an important role in determining the best database design for an enterprise.

# Summary of ER (Contd.)

- ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
  - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, and whether or not to use aggregation.
- Ensuring good database design: resulting relational schema should be analyzed and refined further. FD information and normalization techniques are especially useful.